

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.callbacks import EarlyStopping
```

```
In [9]: # Load datasets
train_data = pd.read_csv('bank-full_train.csv')
test_data = pd.read_csv('bank-full_test.csv')
```

```
In [13]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler

# EDA - Exploratory Data Analysis
# EDA - Exploratory Data Analysis
def eda_visualizations(data):
    # Handle missing and invalid data
    data.replace([np.inf, -np.inf], np.nan, inplace=True)
    data.dropna(inplace=True) # Drop rows with NaN values

    # Age distribution
    plt.figure(figsize=(10, 6))
    sns.histplot(data['age'], kde=True)
    plt.title("Age Distribution")
    plt.savefig('eda_age_distribution.png')

    # Job distribution
    if 'job' in data.columns:
        plt.figure(figsize=(10, 6))
        sns.countplot(x='job', data=data, order=data['job'].value_counts().index)
        plt.xticks(rotation=45)
        plt.title("Job Type Distribution")
        plt.savefig('eda_job_distribution.png')
```

```
# Balance by subscription
if 'y' in data.columns:
    plt.figure(figsize=(10, 6))
    sns.boxplot(x='y', y='balance', data=data)
    plt.title("Balance by Subscription")
    plt.savefig('eda_balance_by_subscription.png')

# Correlation heatmap
numeric_data = data.select_dtypes(include=[np.number]) # Select only numeric columns
if numeric_data.shape[1] > 1: # Check if there are enough numeric columns
    plt.figure(figsize=(10, 6))
    corr = numeric_data.corr() # Compute correlation on numeric data only
    sns.heatmap(corr, annot=True, cmap="coolwarm")
    plt.title("Correlation Heatmap")
    plt.savefig('eda_correlation_heatmap.png')
else:
    print("Not enough numeric columns for correlation heatmap.")

# Preprocess data
def preprocess_data(data, is_train=True):
    # Handle missing and invalid data
    data.replace([np.inf, -np.inf], np.nan, inplace=True)
    data.dropna(inplace=True)

    # Encode categorical columns
    label_encoder = LabelEncoder()
    categorical_cols = ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'outcome']
    for col in categorical_cols:
        if col in data.columns:
            data[col] = label_encoder.fit_transform(data[col])

    # Encode target column (only for training data)
    if 'y' in data.columns:
        data['y'] = label_encoder.fit_transform(data['y'])

    # Separate features and target
    features = data.drop(columns=['y', 'ID']) if is_train else data.drop(columns=['ID'])

    # Scale numeric features
    scaler = StandardScaler()
    features = pd.DataFrame(scaler.fit_transform(features), columns=features.columns)
```

```
    return features, data['y'] if 'y' in data.columns else None

# Apply EDA and Preprocessing
print("Checking data type for train_data:")
print(type(train_data)) # Debugging step to ensure train_data is a DataFrame


eda_visualizations(train_data)

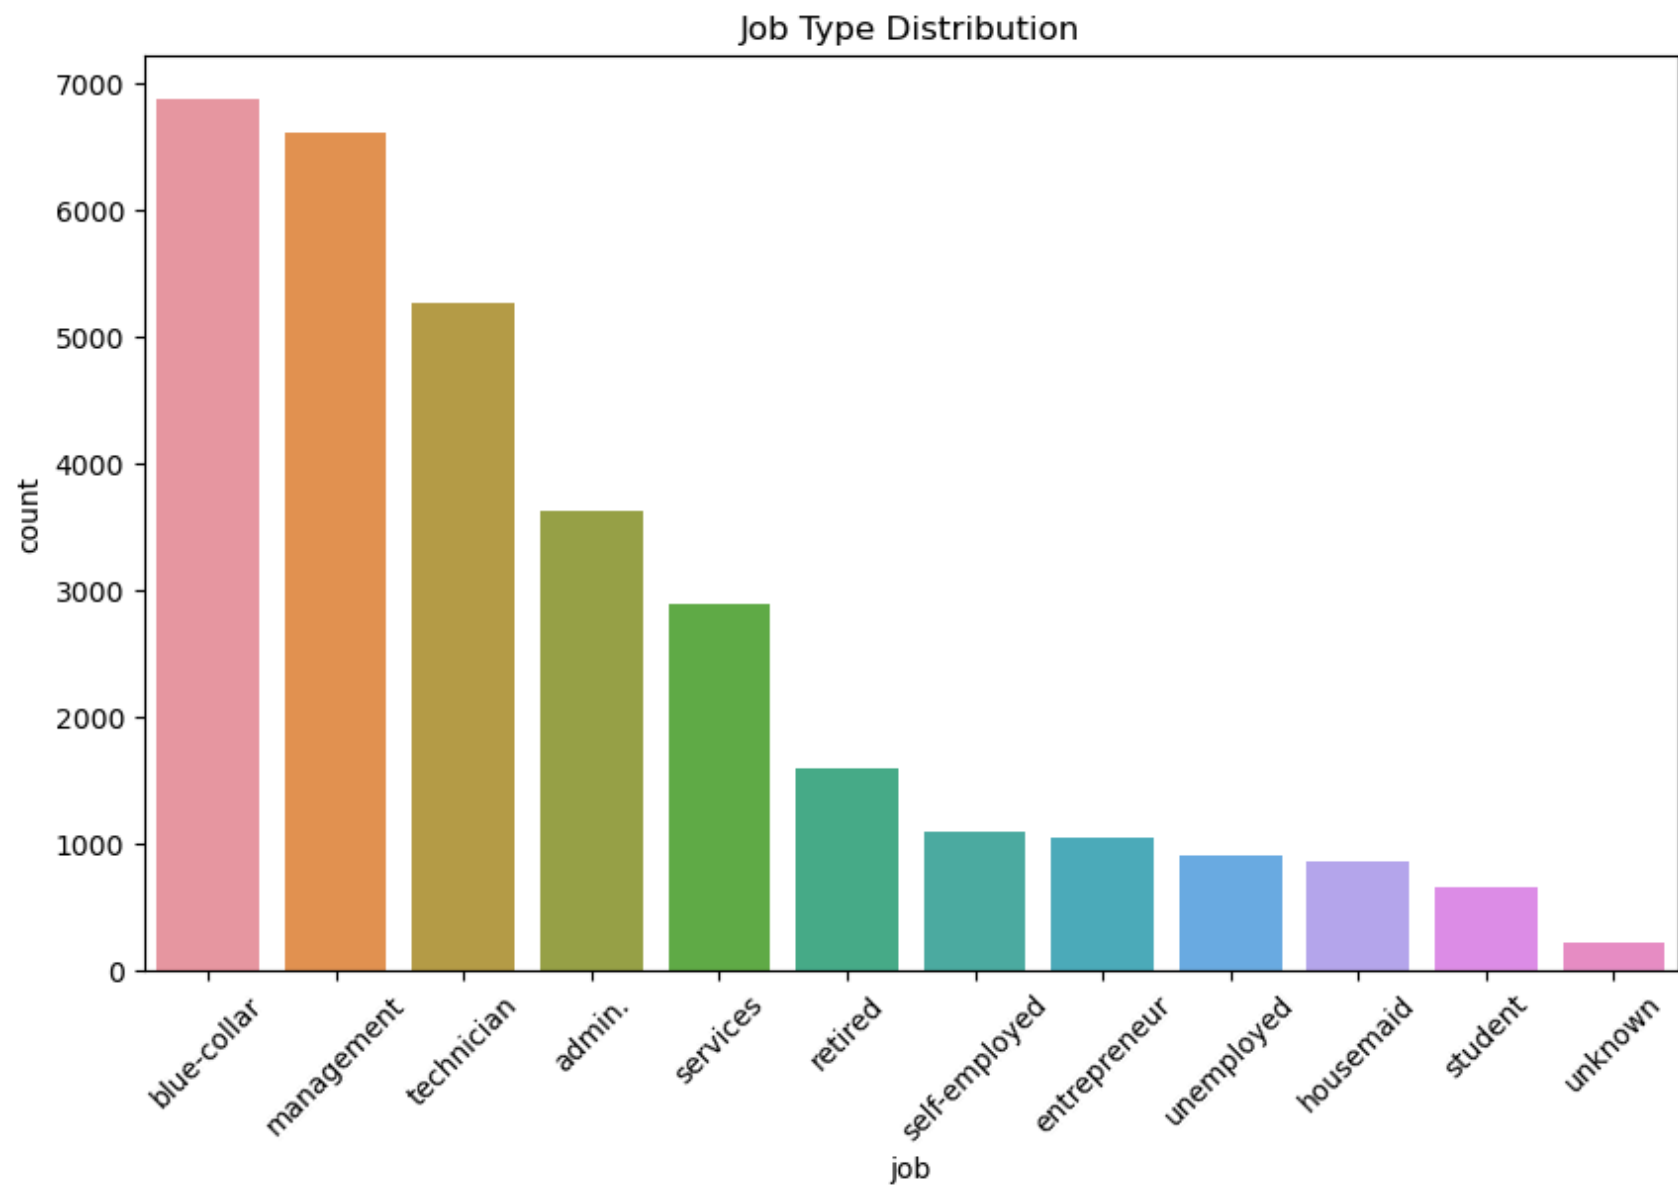
X_train, y_train = preprocess_data(train_data)
X_test, _ = preprocess_data(test_data, is_train=False)
```

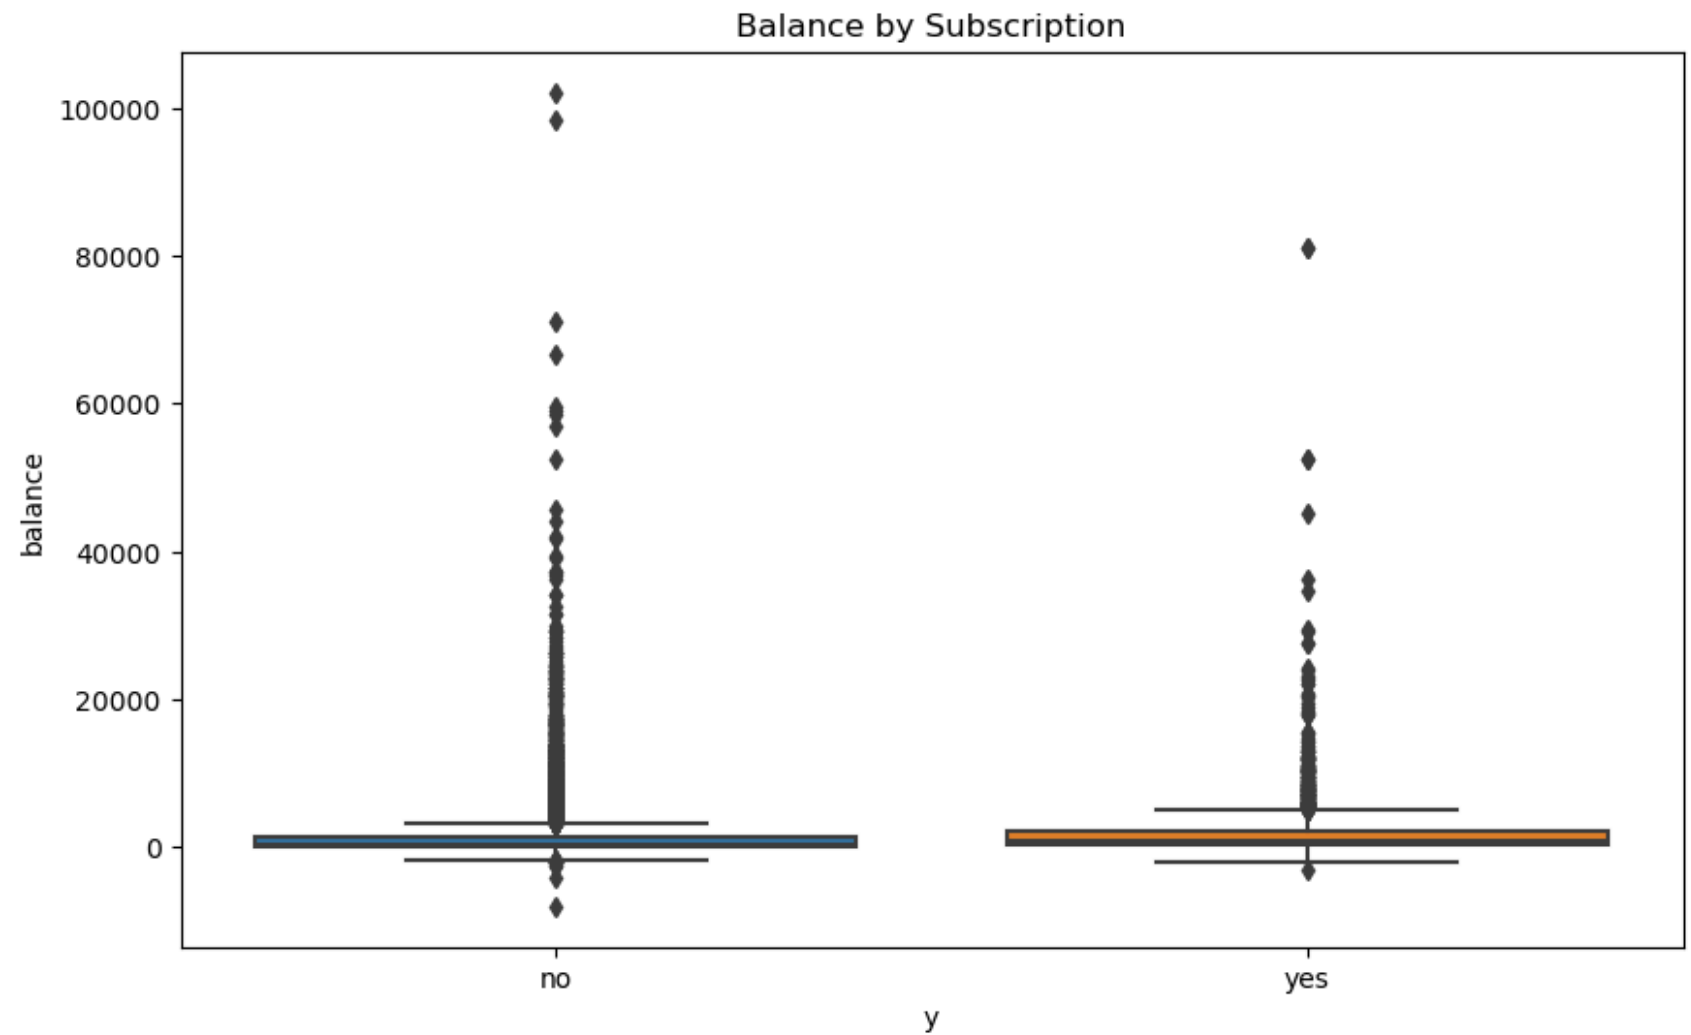
Checking data type for train\_data:

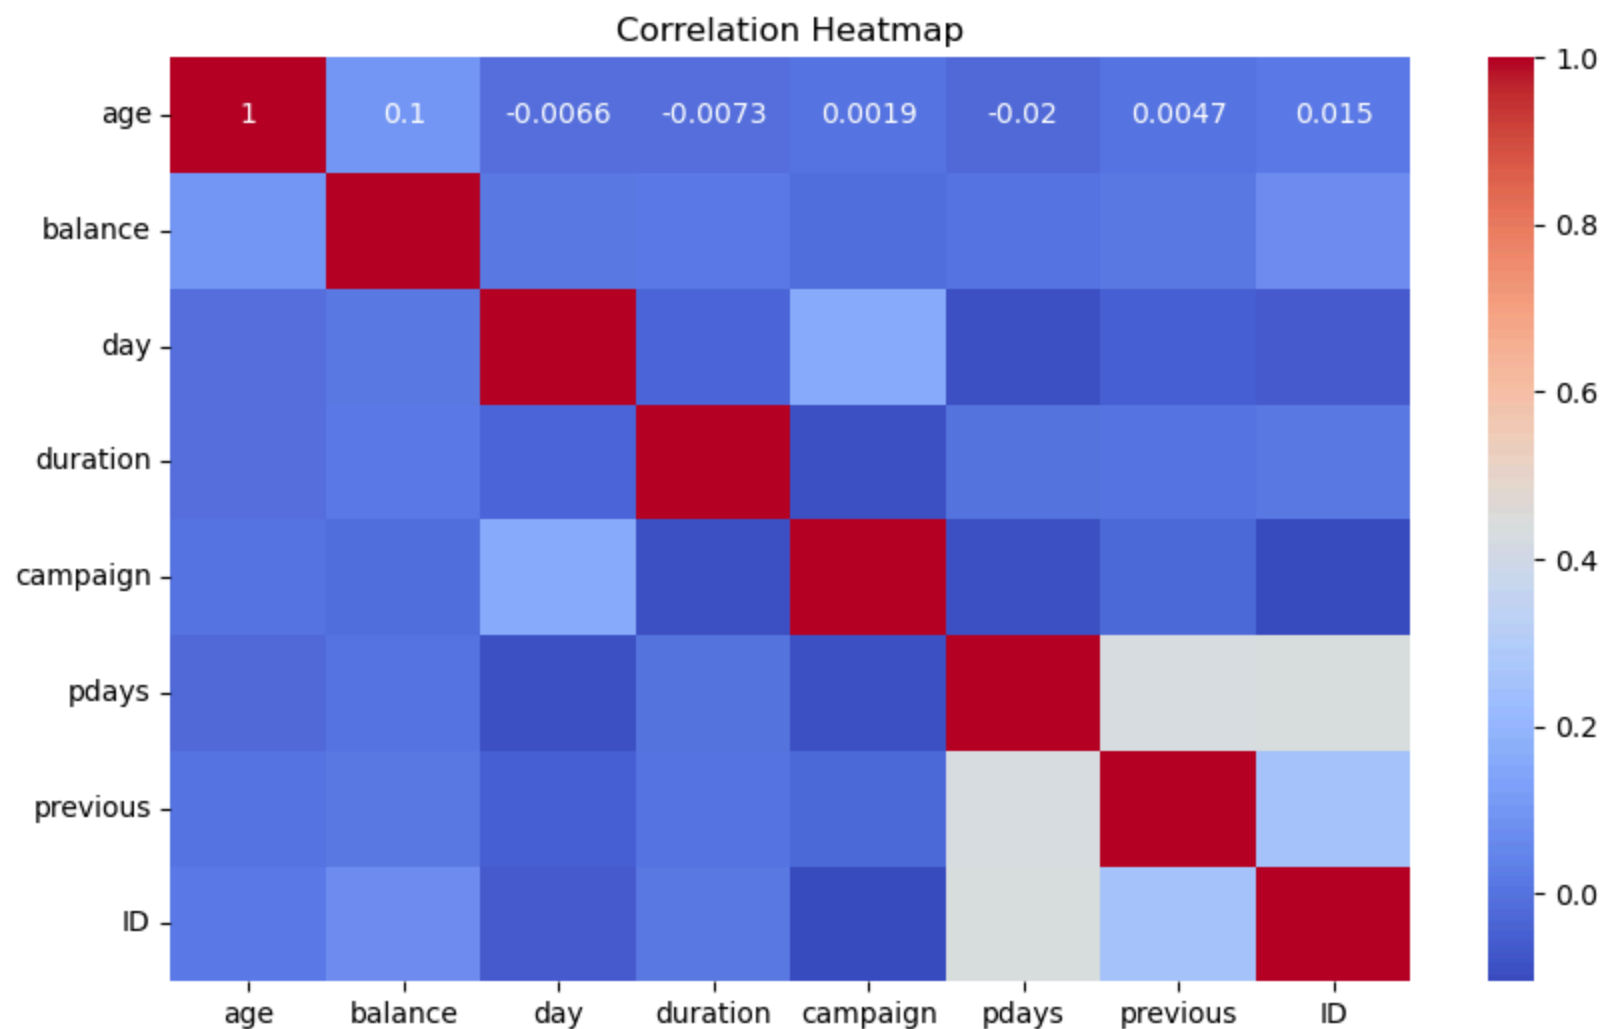
<class 'pandas.core.frame.DataFrame'>

C:\Users\vpark\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):









```
In [14]: # Neural Network Model
def create_nn_model(input_shape):
    model = Sequential([
        Dense(64, activation='relu', input_shape=(input_shape,)),
        Dense(32, activation='relu'),
        Dense(16, activation='relu'),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
    return model
```

```
# Model Training
model = create_nn_model(X_train.shape[1])
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
history = model.fit(X_train, y_train, validation_split=0.2, epochs=50, batch_size=32, callbacks=[early_stopping])

# Save model and training history plots
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.title('Model Loss Over Epochs')
plt.savefig('nn_model_loss.png')

# Predictions on Test Data
y_test_pred = model.predict(X_test)
y_test_pred_class = (y_test_pred > 0.5).astype(int)

# Model Performance Evaluation on Training Data
y_train_pred = model.predict(X_train)
y_train_pred_class = (y_train_pred > 0.5).astype(int)
fpr, tpr, thresholds = roc_curve(y_train, y_train_pred)
roc_auc = roc_auc_score(y_train, y_train_pred)

# Plot ROC Curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Training Data')
plt.legend(loc="lower right")
plt.savefig('nn_model_roc_curve.png')

# Confusion Matrix Visualization
conf_matrix = confusion_matrix(y_train, y_train_pred_class)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", cbar=False)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Training Data")
plt.savefig('nn_model_confusion_matrix.png')
```



```
# Output predictions to CSV for review
test_data['y_pred'] = y_test_pred_class
test_data[['ID', 'y_pred']].to_csv('bank_test_predictions.csv', index=False)
```

C:\Users\vpark\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/50

792/792 ————— 4s 3ms/step - accuracy: 0.8455 - loss: 0.3551 - val\_accuracy: 0.8973 - val\_loss: 0.2434

Epoch 2/50

792/792 ————— 2s 2ms/step - accuracy: 0.8991 - loss: 0.2359 - val\_accuracy: 0.8962 - val\_loss: 0.2348

Epoch 3/50

792/792 ————— 2s 2ms/step - accuracy: 0.8998 - loss: 0.2311 - val\_accuracy: 0.8987 - val\_loss: 0.2322

Epoch 4/50

792/792 ————— 2s 2ms/step - accuracy: 0.9041 - loss: 0.2262 - val\_accuracy: 0.8967 - val\_loss: 0.2308

Epoch 5/50

792/792 ————— 2s 2ms/step - accuracy: 0.9040 - loss: 0.2190 - val\_accuracy: 0.8979 - val\_loss: 0.2297

Epoch 6/50

792/792 ————— 2s 2ms/step - accuracy: 0.9044 - loss: 0.2230 - val\_accuracy: 0.9005 - val\_loss: 0.2327

Epoch 7/50

792/792 ————— 2s 2ms/step - accuracy: 0.9037 - loss: 0.2230 - val\_accuracy: 0.8972 - val\_loss: 0.2269

Epoch 8/50

792/792 ————— 2s 2ms/step - accuracy: 0.9079 - loss: 0.2188 - val\_accuracy: 0.8962 - val\_loss: 0.2352

Epoch 9/50

792/792 ————— 1s 2ms/step - accuracy: 0.9022 - loss: 0.2209 - val\_accuracy: 0.8973 - val\_loss: 0.2300

Epoch 10/50

792/792 ————— 2s 2ms/step - accuracy: 0.9053 - loss: 0.2141 - val\_accuracy: 0.8959 - val\_loss: 0.2294

Epoch 11/50


792/792 ————— 2s 2ms/step - accuracy: 0.9095 - loss: 0.2050 - val\_accuracy: 0.8992 - val\_loss: 0.2271


Epoch 12/50


792/792 ————— 1s 2ms/step - accuracy: 0.9076 - loss: 0.2125 - val\_accuracy: 0.8956 - val\_loss: 0.2346

424/424 ————— 0s 838us/step

989/989 ————— 1s 890us/step







In [ ]: